

Weekly Report

CS3500: Operating Systems

Visualisation Tool for Process Scheduling



Computer Science and Engineering
Indian Institute of Technology Madras
Jul - Nov 2024

Under the supervision
of
Prof. Janakiram D
submitted by
Team 8

- Anjali Samudrala (CS22B046)
- Chaitanya Sai Teja G (CS22B036)
- Jwala Likitha Reddy M (CS22B078)
- Karthikeya P (CS22B026)
- Naveen Koushik Reddy E (CS22B006)
- Navya Sree B (CS22B045)
- Rushi Babu G (CS22B040)
- Sravya Rangu (CS22B044)
- Yashwanth Sai P (CS22B002)
- Yaswanth Sai V (CS22B043)

1 Frontend Team

1.1 Tasks

- The tasks for this week is to identify the components that are required for the frontend. This includes the following.
 - Coming up with a frontend design and layout.
 - Ideation on what all different plots can be displayed.
 - Informing the backend team about the requirements of various informations.

1.2 Progress

1.2.1 Frontend Design and Basic implementation

-

1.2.2 Ideas for Plots and information from backend

- **Processes Running on the System:**
 -
- **Gantt Chart of various Processes**
 -
- **Status wise distribution of processes**
 -
- **CPU Utilization**
 -

2 Backend Team

2.1 Extracting Process Information

2.1.1 Command: `ps -eo pid,lstart,comm`

Purpose: This command displays the details of all currently running processes with the following information:

- **PID:** The Process ID, which uniquely identifies each process.
- **Lstart:** The start time of the process.
- **Comm:** The command or program name that initiated the process.

This command is useful for tracking when a process was started and understanding its origin.

2.1.2 Command: `ps -eo pid,etime,comm`

Purpose: This command displays the following details for all running processes:

- **Etimes:** The elapsed time since the process started, which tells us how long the process has been running.

This information is useful for understanding how long a process has been active since it started.

2.1.3 Command: `cat /proc/[pid]/stat`

Purpose: This command is useful for retrieving detailed information about a specific process. It provides the following data:

- **State:** The current state of the process (e.g., running, sleeping, etc.).
- **Utime:** The total time the process has been running in user mode. This value represents the amount of CPU time consumed by the process in user mode.
- **Stime:** The total time the process has been running in kernel mode. This value represents the CPU time spent by the process in kernel space.

This command is valuable for getting low-level process statistics and understanding the process's CPU usage in both user and kernel modes.

2.1.4 Command: `ps aux`

Purpose: This command provides more detailed information about all running processes, including:

- **State:** The current state of the process.
- **Start time:** The start time of the process.
- **Total CPU time:** The total CPU time the process has consumed.

This command is useful for monitoring various aspects of processes and understanding their current state and resource usage.

2.1.5 Command: `pidstat`

Purpose: The `pidstat` command is used to gather detailed statistics for individual processes. It can provide more granular information about resource utilization, including CPU usage, memory usage, and more, for each process.

This tool is helpful for detailed performance analysis and understanding how specific processes affect system resources.

2.1.6 Command: `ps -p pid`

Purpose: To know whether the process ended or not.

- If it shows some data, it is still running.
- If no data is shown, the process ended.

Consider a specific time duration. During this time, we generate a list of all processes currently running. After this duration has passed, we check the list of running processes again. If we find that any process from the initial list is no longer running, it indicates that this process has completed or ended.

2.2 Linking Backend and Frontend

The objective is to create a backend server that fetches real-time system process statistics using the `pidstat` command and streams the data to a frontend via WebSockets. The application uses Flask as the web framework and Flask-SocketIO to establish real-time communication between the backend and frontend.

2.2.1 Tools and Libraries Used

- **Flask**: A lightweight web framework for Python that simplifies the development of web applications.
- **Flask-SocketIO**: An extension for Flask that enables real-time communication between the server and the client using WebSockets.
- **subprocess**: A Python module used to run external commands. In this case, it is used to execute the `pidstat` command, which collects CPU statistics for processes running on the system.
- **re (Regular Expressions)**: A Python module for matching patterns in strings. It is used to parse the output of the `pidstat` command.
- **threading**: A Python module used to create background threads. In this project, it allows the data-fetching process to run concurrently with the main server.

2.2.2 Architecture

- **Backend (Flask Server with SocketIO)** : The backend is responsible for fetching real-time process statistics using the `pidstat` command, parsing the output, and sending the data to the frontend using WebSockets.
- **Real-Time Data Fetching** :
 - The `pidstat` command is used to gather CPU statistics for processes every second. The output contains several fields, including process ID (PID), user and system CPU usage, and the process command.
 - A background thread is created to run the `pidstat` command continuously, fetching data at specified intervals (e.g., every 2 seconds).
 - The data is parsed using a regular expression, and relevant statistics are extracted and formatted into a dictionary.
- **Real-Time Communication** :
 - `Flask-SocketIO` is used to emit the parsed data to the frontend in real time. This allows the frontend to display the latest statistics as they are gathered by the backend.
 - The data is emitted as a WebSocket event (`'pidstat_data'`), making it available for frontend visualization.

2.3 Process Migration

2.3.1 Enabling tracking to log the process migration

- Navigate to the tracing directory.
`bashcd /sys/kernel/debug/tracing`
- Enable "sched_migrate_task" that allows to track and log the migration of tasks.
`bashecho 1 — sudo tee events/sched/schedmigratetask/enableStarttracingtheevents.`
`bashecho 1 — sudo tee tracingon`
- Wait for some time and let tracer log some migrations.
- Stop the tracing.
`bashecho 0 — sudo tee tracingonCheckthetracefile.`
`bashcat trace`

2.3.2 Logged Data

- An example of the logged data is shown below. `text ChromechildIOT-33594[003]d..2.27300.882725 : schedmigratetask : comm = Compositorpid = 32759prio = 120origcpu = 0destcpu = 3`
- In the above log entry, Chrome with PID 33594 initiated the migration of the Compositor process with PID 32759 from CPU 0 to CPU 3.
- `prio = 120` is the priority of the process. Lower prio means higher priority.
- `27300.882725` is the timestamp of the event. (migration).

2.4 Completely Fair Scheduler(CFS) Understanding

- As we are implementing this project on Linux OS, an understanding of its scheduler is beneficial.
- It assigns priority to a process based on its niceness value which in turn is based on the vruntime of the process.
- The process with the least vruntime is given the highest priority. vruntime (virtual runtime) tracks the time a process has spent on the CPU, adjusted for its priority.